Data warehousing of European parliament voting data and Unsupervised Learning

By **Franciszek Świderski**

B. Sc. International Politics and Government, Bocconi University

Under the supervision of

Omiros Papaspilioupoulos

Full Proffesor at Department of Decision Sciences, Bocconi University

For Institute for European Policymaking at Bocconi University

October 2024

Abstract

Abstract goes here

Contents

Chapter 1

Introduction

This is an introduction to the work, I do some intorducting

1.1 Dataset Structure and Initial Challenges

The dataset comprises three primary components:

- MepInfo: This dataset contains information on Members of the European Parliament (MEPs).

- Votes: This dataset records the votes cast by MEPs on specific pieces of legislation.

- Votings: This dataset describes the legislative items that were voted upon.

FullName	EPG	1	2	3	...	18674
Charles GOERENS	RE	5	1	1	...	2
Bill NEWTON DUNN	RE	2	1	1	...	0
Constanze KREHL	S&D	3	1	1	...	0
Coelho COELHO	EPP	0	0	0	...	3
Bernd LANGE	S&D	3	1	1	...	2
...	...	...	...	...	...	...
Henk Ormel	EPP	0	0	0	...	1

Figure 1.1: Data from European Parliament 9 formatted for roll-call scaling

The Votes dataset alone includes approximately 42,800 roll-calls, with 800-900 MEPs participating in each legislature, resulting in an estimated 40 million data points.
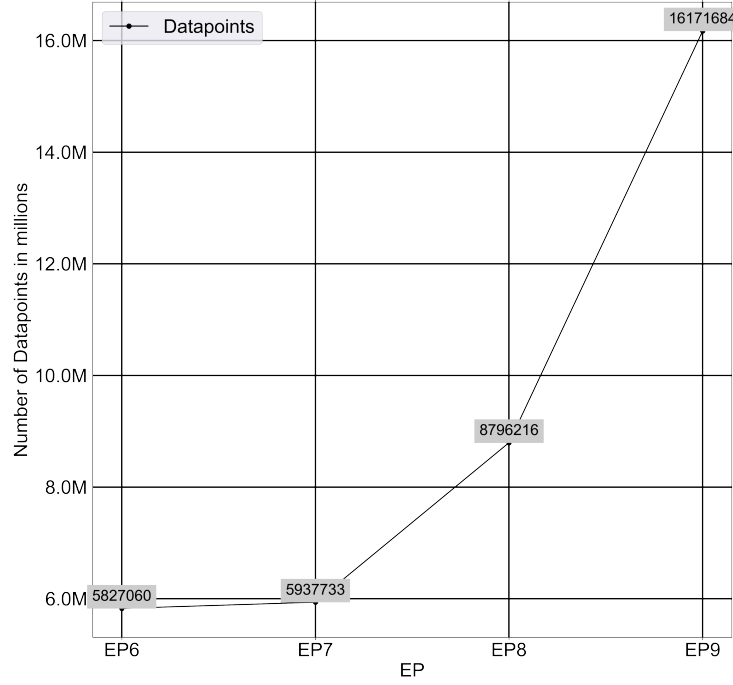
Figure 1.2: Number of datapoints by European Parliament (EP)

The size of the dataset makes for a serious practical challenge in terms of storage and processing, not to mention the problems with formatting and integrity of the dataset. Upon initial examination, several issues with the legacy data became apparent:

- Missing Variables: The MepInfo dataset lacked several critical variables, including gender and age.

- Inconsistencies and Missing Observations: Some variables such as Party affiliation and European Parliament Group (EPG) affiliation were either inconsistent or incomplete.

- Non-Standard MEP IDs in EP7: MEP IDs in the 7th Parliament did not conform to the standard European Parliament API format. Instead of unique identifiers, incrementing integers were used.

- Naming Conventions: MEP names and surnames were not consistent with those in the European Parliament API.

- Incorrect and Inconsistent Encoding of Missing Data: Missing data was inconsistently encoded, with symbols like - or . being used in place of null values.

- Inconsistent Datetime Formatting: Datetime values were formatted inconsistently across the dataset, making them difficult to import and process.

- Non-Standard Encoding of Binary Categorical Variables: For example, the binary categorical variable " Vote," which indicates whether a vote passed or not, was encoded as + and - instead of the standard 1/0 used elsewhere in the dataset.

Additionally, the dataset was stored in a long format tailored to the specific analytical requirements of the original VoteWatch team. While this structure was suited to their immediate needs, it posed challenges for long-term usability, warehousing, and broader analytical purposes.

Chapter 2

Data Cleaning Process

2.1 Supplementing and Correcting the MepInfo Dataset

The initial challenge involved supplementing and correcting the MepInfo dataset to fill in missing variables and rectify inconsistencies. The European Parliament's API played a crucial role in this process. An API (Application Programming Interface) is a set of defined rules and protocols that allows one software application to interact with another. In the context of scraping data, an API serves as an intermediary that lets users request and retrieve specific data from a website or service in a structured format, typically JSON or XML, without the need to manually scrape the HTML content of a webpage. Using an API for data scraping is often more efficient and reliable than traditional web scraping, as it provides direct access to the desired data, reducing the risk of encountering issues like changes in webpage structure or content restrictions. This API provides endpoints, such as '/meps', which return JSON data containing basic MEP information, including a unique identifier (MepId). By utilizing this API, we were able to standardize the names and identifiers of MEPs across most legislatures by joining the tables by MepId.

However, the 7th Parliament (EP7) presented a unique challenge. The MEP IDs in this legislature were not unique and were instead represented by incrementing integers. To address this issue, we employed the 'fuzzywuzzy' Python package, which uses the Levenshtein distance algorithm to calculate the similarity between strings. This allowed for making an approximate match of full names from the original dataset to the 'sortLabel' field in the API data, providing the correct MEP IDs. Manual verification and correction of edge cases were necessary to ensure accuracy.

Despite these efforts, some data remained incomplete, particularly regarding MEPs who changed parties or EPGs during their tenure, as well as demographic data such as gender and birth dates. These gaps required further manual supplementation, which is discussed in detail in the Automation section.

2.2 Addressing Inconsistencies in the Votings Dataset

The Votings dataset required significant work to address encoding inconsistencies and improperly formatted datetime values. The first issue was relatively easy to reconcile using renaming dictionaries to replace inconsistent encodings in some variables.

However, the latter issue proved particularly problematic due to the original data being gathered in Excel. Manual formatting likely led to inconsistent datetime values formatting. To standardize the format, manual correction of the dates was necessary before applying the Pandas 'to datetime' function.

This step ensured that the datetime values were correctly parsed and made the data suitable for further analysis. Once the initial cleaning was complete, a new challenge arose. In 2022, the original VoteWatch team altered their data gathering methodology. Consequently, data from Parliament 9, covering the period until 2022, was consistent with earlier practices. However, data collected from 2022 to March 2024 exhibited several discrepancies. The order of VoteIds had changed, some columns particularly those in the Votings dataset were missing, and special characters were corrupted due to encoding issues. To resolve these issues, we employed a multi-step approach. First, we used a dictionary to replace the corrupted characters systematically. Next, we reverse-engineered the missing Votings columns (such as finalVote, a binary variable of whether a vote is a last one in a section) from the available data. Despite the complexity of these tasks, they were necessary to restore the integrity and consistency of the dataset, ensuring that it could be seamlessly integrated with the existing data from earlier legislative periods.

2.3 Data Restructuring for Long-Term Usability

With the Votes and MepInfo datasets cleaned and supplemented, the next step was to restructure the data into a format suitable for long-term storage, analysis, and future updates. The data was initially stored in a long format, where each row represented an observation, and columns represented variables. To facilitate analysis, we separated the MepInfo and Votes datasets into distinct tables:

- MepInfo table contained all variables describing the MEPs, with 'MepId' serving as the primary key.

- The Votes table used a composite primary key consisting of 'MepId' and 'VoteId', a unique identifier for each voting event. This table also included a column encoding the MEPs' votes.

These tables were linked by primary keys, allowing for efficient querying and data retrieval. For example, using the 'MepId', one could easily retrieve all votes

cast by a particular MEP, and by further referencing the 'VoteId', additional context from the Votings table could be added.

To support the development of the European Parliament Vote Monitor website, the voting data was additionally exported in '.csv' format, organized by month and year. This format was chosen to optimize data storage and accessibility, ensuring that the data could be easily updated and queried. Additionally, for specialized analytical purposes, the Votes data was also retained in its original matrix format, now enhanced with the newly cleaned and supplemented variables.

Chapter 3

Automation of Data Gathering

3.1 Overview of the European Parliament API

With the historical data cleaned and restructured, the next phase of the project focused on automating the future data gathering process. Consistency and reliability were paramount in this process. The previous VoteWatch team relied on scraping XML files from the human-readable official minutes on the European Parliament's website, coupled with downloading MepInfo from the API. This approach was fraught with challenges, including the potential for inconsistencies in the scraped data, lack of unique identifiers, and the inherent unreliability of scraping human-readable content.

To address these issues, we transitioned to using the European Parliament's open API, which provides direct access to structured data in a reliable and consistent manner. The API is publicly accessible, meaning no API key is required, and it allows users to specify parameters in the URI to retrieve data from various endpoints.

3.2 MepInfo Data Collection Automation

The '/meps' endpoint of the API, which provides a list of MEPs along with their unique identifiers, was the starting point for automating the MepInfo data collection. However, the data returned by this endpoint was not exhaustive. To gather more detailed information such as MEPs' gender, age, and political affiliations additional API calls were required to other endpoints within the 'meps' group.

For example, to retrieve data on MEPs' party memberships and affiliations with European Political Groups (EPGs) and National Parties, separate API calls were made specyfing the MepId . The resulting data was then filtered and

joined with the MepInfo dataset to create a comprehensive record of each MEP s political affiliations and demographic information.

One key challenge in this process was optimizing the API calls to minimize the load on the server. Each MEP required an individual API call to gather detailed data, which could result in approximately 850 GET requests per scraping session. To mitigate this, the metadata was included in the same call to avoid the need for additional requests.

Another consideration was the use of non-human-readable IDs for certain values, such as the organization field, which required further API calls to retrieve the corresponding labels from the 'corporate-bodies' endpoint. For example org/1537 was an identifier for GUE/NGL EU Political Group. These labels were essential for understanding the data and were subsequently joined with the MEPs' political affiliation data.

3.3   Votings Data Collection Automation

The 'meetings' endpoint of the API was utilized to automate the collection of data on plenary sessions and the decisions made during these sessions. An initial API call returned a list of all plenary sessions that took place in a given year. This list was then filtered by month to identify the specific sessions relevant to the current data collection period.

For each identified session, further API calls were made to retrieve detailed voting data, such as the votes cast by each individual member and all of the supplementary data that was previously found in the Votings database. This data was then merged with the previously gathered MepInfo and Votes data, ensuring that all relevant information was captured and linked across the collection.

In order to automatically gather the data in the future, working with a Bocconi IT Department team, we created a cloud-based solution that is deployed and triggers every month, gathering and exporting the data to the .csv format, as well as to a relational SQL database.

Chapter 4

Ideal Points Estimation - Overview of Current Methods

One of the most prominent uses of voting data in political science research is estimating ideal points of legislators on the ideological plane. The early revolutional work of Poole and Rosenthal (1985) set the precedent for using those types of algorithms - specifically NOMINATE - Nominal Three Step Estimation on larger datasets than previously possible. Ever since, ideal points estimation was developed more and more, with improved versions of the original algorithm (such as D-NOMINATE, and later W-NOMINATE), to extending its implementation beyond the original usecase (scaling the US parliament), analysing the European Parliament and various national parliaments. The European parliament specifically has had a great deal of coverage using the WNOMINATE method, which has been the standard for the past 20 years.

4.1 WNOMINATE

4.1.1 Algorithm

W-NOMINATE (Weighted NOMINATE) - an improved version of NOMINATE, used for estimating legislator's ideal points relying on their roll-call voting records. The underlying model of NOMINATE, and by extention, W-NOMINATE is a spatial voting model, where legislators and the policies are represented in a multidimensional space. In this implementation, a legislator has an ideal point in this space, that represents their ideology and policy preference, and each vote is a set of two alternative points - one when the policy is implemented (the vote passes), and another when the policy is rejected (the vote fails). W-NOMINATE assumes that legislators vote for an outcome that is closest to their ideal point, with some random error to account for unpredictability.

The utility a legislator derives from a vote is modeled as a function of the Euclidean distance between their ideal point and the vote outcome locations. Legislators aim to maximize their utility by voting for the outcome that minimizes the distance to their ideal point. The model uses the following utility function:

$$U_{ijy} = \beta \exp\left[-\sum_{k=1}^{s} w_k^2 (x_{ik} - z_{jy})^2 / 2\right]$$

where:

- x_{ik} is the ideal point of legislator i on dimension k,

- z_{jy} is the position of the "yea" outcome for vote j,

- β controls the ratio of deterministic to stochastic utility,

- w_k is a weight for each dimension.

The likelihood of voting "yea" is determined by comparing the utilities of the two possible outcomes (yea or nay). W-NOMINATE estimates the legislators' ideal points and vote outcome locations iteratively, using maximum likelihood estimation. The process continues until the estimates stabilize and a convergence threshold is met. Typically, one or two dimensions are used in most applications, with the first dimension often capturing the left-right ideological spectrum.

4.1.2 Initialization

The initialization of W-NOMINATE requires setting initial values for the legislators' ideal points and vote outcome positions. These are assigned by estimating the positions of the legislators by calculating an agreement matrix. It is calculated by assigning values (in this case, they were most likely chosen heuristically) to each vote, Yea is 1, Nay is 6 and Missing (abstentions, etc.) are 9. Then the program iterates over each vote in each pair of legislators, calculating the average distances between their votes and storing it in a square matrix. Then the eigenvectors of this matrix are used as starting points for the algorithm to iterate

The key parameter β is usually set to an initial value of 15, indicating that the deterministic component of utility strongly outweighs the stochastic component.

Additionally, a "lop" threshold is set to determine which votes are included in the analysis. This threshold excludes votes where too many legislators vote in the same way, as these votes provide little information about ideological differences. The default value is 0.025, meaning that votes where fewer than 2.5% of legislators voted in the minority are excluded.

4.1.3 Implementation

W-NOMINATE is implemented in the `wnominate` package for `R` , which builds upon earlier Fortran-based implementations. The package simplifies data input and provides tools for estimating ideal points and visualizing results. Roll-call data is typically stored in a `rollcall` object from the `pscl` package. The iterative maximum likelihood estimation process continues until the ideal points and vote outcomes converge, defined as a correlation of 0.99 between successive iterations.

The `wnominate` package also includes an option for bootstrapping to estimate standard errors. Users can specify the number of bootstrap trials, with the default being no bootstrapping. Visualizations, such as plots of ideal points and vote outcome locations, are also provided to facilitate interpretation of the results.

While the package has been the standard in political sciences for the past 20 years, the software implementation has its' limitations. Lack of multi-threading support does not take advantage of the computing power of new machines effectively and the algorithm is prohibitively slow when working with larger datasets, such as the European Parliament 9 votes.

4.2 emIRT

4.2.1 Algorithm

In order to circumvent the limitations of the traditional methods Imai, Lo, and Olmsted (2016) developed a new algorithm, emIRT - it is particularly efficient in estimating the ideal points in large datasets - where traditional methods, such as Markov Chain Monte Carlo (MCMC) or the NOMINATE procedure , are computationally expensive. emIRT uses the Expectation-Maximization (EM) framework, which is faster and more scalable.

At the core of emIRT is an item response theory (IRT) model, where legislators' votes are treated as indicators of latent ideological preferences (ideal points). The algorithm estimates ideal points by iterating between two steps:

- **E-step** : The algorithm computes the expected value of the latent variables (the legislators' underlying voting propensities) given the current estimates of the ideal points and vote parameters.

- **M-step** : It updates the estimates of the ideal points and discrimination parameters by maximizing the expected log-likelihood calculated in the E-step.

The process repeats until convergence. Unlike MCMC, EM is deterministic, converging faster and without the need for sampling. emIRT can handle binary, ordinal, and continuous voting data, making it highly flexible. It can also extend to dynamic and hierarchical models, where ideal points change over time or across groups.

4.2.2 Initialization

emIRT requires random initialization for the ideal points and the discrimination parameters.

4.2.3 Implementation

The emIRT algorithm is implemented in the `emIRT` package for `R` , which is designed to handle large-scale datasets efficiently. The package includes several models, including:

- **Binary IRT models**: For binary voting data,

- **Ordinal IRT models**: For ordinal voting data,

- **Dynamic IRT models**: For time-varying ideal points,

- **Hierarchical IRT models**: For nested data structures.

The algorithm performs Expectation-Maximization to estimate the ideal points and vote parameters. It can also estimate standard errors using a parametric bootstrap, simulating datasets based on the estimated model and then re-estimating the parameters to compute uncertainty. The package handles missing data in roll-call voting datasets by treating abstentions as latent variables to be predicted from the observed data. This ensures that the ideal points reflect complete voting behavior.

Chapter 5

Ideal points estimation - using the software (very much working name)

Having understood, cleaned and restructured the data, the natural next step was using it for research. However, before going further, we needed to verify the integrity of the data - whether the cleaning process impacted the dataset in a meaningful way that would distort the results. In order to achieve this, we set out to recreate the workflow that Hix and Noury (all the years) used to reproduce their results. If the results were comparable to theirs, that would mean the dataset is functionally the same, as well as allowing us to use the workflow for further research.

5.1 Reproduction of results in Hix Noury 2009

The first step in the process of veryfing the integrity of the data was to establish a reference point of comparison. The data we are working with spans from EP6 to EP9. The publications where Hix Noury explicitly state the use of NOMINATE to create the two-dimensional representations of ideology in the European Parliament cover the range of EP1 to EP6 (later publications, like Hix, Noury Rolland (2018) do not explicitly state the method of roll-call scaling and the results point to use of another method). This meant that the workflow reproduction is feasible in EP6, so we decided to pursue this direction.

5.1.1 Workflow reproduction

Having chosen the specific datset to produce the results from, we begun investigating the `wnominate` package for `R` described in Chapter 4. The documentation of the package in CRAN (The Comprehensive R Archive Network - the pack-

age manager for R) lays out the steps for scaling an arbitrary (e.g. not from the US parliament) dataset of roll-call votes. The first step is importing the dataset and converting it to a `rollcall` object from the `pscl` package. This object store the votes and decsriptive data about legistators in accordance with the `wnominate` requirements. It recodes the votes to only contain 4 categories - 'yea' , 'nay' , 'missing' and 'not in legislature'. As previously mentioned, the votes are stored in the long format matrix - specific votes are columns, and each row is a legislator. From this point on, the only step required to initialize the algorithm is choosing the parameters β and 'lop', as well as the ' polarity' of the dataset.


Figure 5.1: Ideal points estimates from HixNoury (2009)

As we do not have access to the code used for the analysis done by Hix and Noury, we have to make assumptions as to the parameters that they chose. β and 'lop' have default values, heuristically assigned by the creators of the software, and researchers usually do not change those. The issue arises with the "polarity" parameter. It requires the researcher using the package to choose a 'conservative' MP and the algorithm will center the result around his final

position - so if the 'polarity' is chosen as (10,10), as it is in our case, `wnominate` will center the result around the 10th MEP in the dataset in dimension 1 and 10th MEP in dimension 2. There is no feasible way to recreate the result exactly without the knowledge of the MEP used as the 'conservative'. However, this parameter does not distort the distribution of the ideal points, just their alignment and rotation - so the final product should be still comparable to the results achieved by Hix Noury. 5.1


Figure 5.2: Recreation of Ideal Points Analysis in European Parliament 6 using WNOMINATE

Comparison of results In order to explicitly compare the results we achieved to the ones of Hix Noury (2009) we would need to have access to the original coordinates produced by W-NOMINATE in their research. This data is unfortunately not available publicly, so the most feasible method of comparison is visual analysis of characteristics of both representations of ideal points. One

can immediately notice that the positions of the main political groups in our map are inverted in relation to the map we are comparing with. This is likely due to a different MEP being chosen as the 'conservative' in the initialization of the algorithm. In order to compare the results more adequately, we decided to show a map with both dimensions inverted, to better match the original results.


Figure 5.3: Recreation of Ideal Points Analysis in European Parliament 6 using WNOMINATE with both dimensions inverted

With maps aligned, several similar characteristics become apparent. The EPGs appear in similar clusters, with a similar alignment as in the original. The Socialists (PES in reproduction) are firmly in the range of -0.2 to 0 in the first dimension in both cases and 1 and 0.25 in the second dimension in the reproduction, but between 1 and 0.6 in the original work. The most likely explanation for this discrepancy is the polarity setting in the second dimension - if we chose a MEP that is closer to some of the Socialists in terms of ideal points

than the one in Hix Noury (2009) as a 'conservative' we would naturally bring the Socialists closer to the middle of the map, distorting the image. That seems to be the most competent explanation of this phenomenon. Again, however, we do not have a way of comparing the results directly and lack the knowledge of the 'conservative' used in the reproduced work.

Similarly to the Socialists, other groups - such as Conservatives (EPP-ED), Liberals (ALDE), Nationalists (UEN) and Grens (Greens-EFA) are in comparable clusters to the previous research. The Radical Left (GUE-NGL), the Anti-Europeans (IND/DEM) and the Independents (NI) are located on opposite spectra of the map, along the borders of the unit sphere (an inherent feature of maps produced by W-NOMINATE are points with maximum distance of 1 from the center). This also reflects accurately the reproduced work.

With such similar results, and without other means of confirming the similarity, we can plausibly establish a successful reproduction of the workflow that was used in Hix Nouury (2009). This results allows us to conclude the integrity of the dataset and methods used.

5.2 Comparison of one-dimensional ideal points methods

Talk about computational limitations of WNOMINATE and how we can use emIRT to get similiar results.

5.2.1 emIRT

Talk about how to get the results, how to use the software etc etc. Show some preliminary graphs etc

5.2.2 WNOMINATE

Same stuff as before.

5.2.3 Comparison

Compare the results and show that they are basically the same, so we can use emIRT in base cases instead of WNOMINATE - and will do so in the future, with more bootstraps etc.

5.3 Initialization research

Establishing they are similar ish - we wanted to run emIRT with WNOMINATE starts and WNOMINATE with emIRT starts

5.3.1 Obtaining the starting values of WNOMIANTE

agreement matrix etc etc

5.3.2 Initializing emIRT with WNOMINATE starts

etc etc couldnt manage WNOMINATE to run

5.3.3 Comparing the results

comparing emIRT normal to not normal

5.4 High-dimensional ideal points

Running WNOMINATE for high d is hard etc etc

5.4.1 Obtaining the ideal points using WNOMINATE

what i had to do etc etc

5.4.2 Interpretation of dimensions

normal interpretations, what can you do maybe more, why more dimensions are usually not used so much etc etc

5.5 Principal Component Analysis

kinda to do Obtaining ideal points estimates using PCA

5.5.1 Comapring with starting points of WNOMINATE

5.5.2 Using Procrustes Rotation for high-dimensional comparison

Chapter 6

Analysing the structure of missing data

6.1 Overview of missing data

6.1.1 Types of missing data in the dataset

6.1.2 Describing the missingess

6.2 Impact of ideology on missingness

6.2.1 Model

6.2.2 Results

Chapter 7

Conclusions